

Jigsaw Toxic Comment Classification

COLX 585 Trends in Computational Linguistics Final Project

Amy Lam; Nilan Saha; Lei Song

1 ABSTRACT

In this project, we fine-tuned a pretrained BERT(Bidirectional Encoder Representations from Transformers) model to detect and classify toxic comments in Wikipedia's talk page edits, which is a dataset collected by Jigsaw and Google and posted on a Kaggle competition. Previous academic papers on toxicity detection with various deep neural network architectures such as CNN and RNN have demonstrated some key challenges of this dataset are severe imbalance of toxic versus non-toxic comments and intentional misspelled words. Drawing from previous work's lessons, we utilized BERT's Word-piece embedding to alleviate out-of-vocabulary problem, and addressed the class imbalance problem effectively by down sampling of non-toxic comments. The fine-tuned model achieved a stable and robust macro-F1 score of around 0.90 on the test split (5,494 comments) of original training data, and a macro-F1 score of 0.788 when applying on the bigger Kaggle test data (153,164 comments), across the 6 binary classifiers on the toxicity categorization labels. While the majority class labeling achieved test-split macro-F1 around 0.965, minority class labeling also performed quite well at a score of 0.837 under scarce input data.

2 INTRODUCTION

There are increasingly more number of people who are willing to use the internet to share their views in many aspects. However, we can find many instances of threat and harassment comments online, so sometimes people are reluctant to share their thoughts and opinions. This actually makes it very difficult to have a free conversation or just in general talk about things. Therefore, more and more platforms are interested in building a friendly and safe environment for their users to communicate with each other. So our goal for this project is to detect toxic comments and classify different kinds of toxicity. Our system is a fine-tuned version of one of the popular pre-trained models(BERT) that works really well for this specific task and gives good multi-label performance.

We have done the Jigsaw Toxic Comments Classification task, in which comments are all in English, which is a text classification task in nature. We have trained a neural model to learn about what words/phrases may be insulting, threatening or hurtful to others. The model takes in a sequence of text, and outputs a binary vector of length 6 which denotes which labels the text can be attributed to among toxic, severe toxic, obscene, threat, insult and identity hate.

The project is primarily socially-motivated. The number of people chatting and commenting and just writing content on the Internet have increased exponentially over the past decade. A lot of people do use toxic language which should be dealt by first detecting if it is toxic and then dealing with it some way like hiding it or just deleting it. The goal of this project is to detect such toxic comments/messages and also classify the type of toxicity.

3 RELATED WORK

A few papers published in recent 3 years have tested different machine learning models on toxic comment classification and they served as our basis for model design. Earlier works before BERT paper came out already achieved remarkable toxicity detection results between 0.80-0.90 F1-score with LSTM and CNN, which tackled the long-range dependency issue of lengthy toxic comments, but toxicity categorization results are less appealing at around 0.65-0.75 F1-score.

In [3], authors explored the use of Support Vector Machine(SVM), Long Short-Term Memory Networks(LSTM), Convolutional Neural Networks(CNN) and Multilayer Perceptrons (MLP) with both word-level and character-level embeddings on the same Kaggle toxic comment classification challenge dataset. The authors achieved an accuracy of 0.889 and a F1 score of 0.886 on the binary classification task with a word-level LSTM model that has 3 layers with 32 output units at each layer. A CNN model with kernel size 3 and dropout ratio 0.2 achieved a F1 score of 0.871. According to their experiments, character-level neural models performed much worse. On a much more fine-grained multiple-label toxic comment classification task, the LSTM model was able to achieve a F1 score of 0.706. The LSTM word-level model surprisingly could detect toxicity despite spelling mistakes in the toxic comment.

In [5], authors point out that main challenges of toxic comment classification include long-range dependencies, intentionally misspelled and idiosyncratic out-of-vocabulary words, class imbalance problem and high variance in data or inconsistency in labeling. Authors applied an ensemble approach, combining strong classifiers such as Logistic Regression, Bi-directional RNN, Bi-directional GRU with Attention layer and CNN with pretrained word embeddings from Glove and sub-word embeddings from FastText. The Bi-directional GRU with Attention model outperformed other models but ensemble approach achieved even higher F1 scores of 0.791 for toxicity detection task on Wikipedia comments and 0.793 on another Twitter dataset. For multi-label comment classification task, authors found that ensembling is especially effective on the sparse classes "threat" and "hate". In the post-model error analysis, the paper pinpointed that remaining major prediction errors came from a few areas such as incorrect original labeling; toxicity without swear words, toxic comments framed as rhetorical questions, subtle metaphors and comparisons that require more real world knowledge/context.

Google published the paper "Attention is all you need" [2] in June 2017 and proved BERT could achieve new state-of-the-art results on multiple downstream NLP tasks by feeding very little finetuning data. Applying BERT to toxicity or offensive comment identification are

more recent attempts and fewer papers are available. Relevant papers showed two approaches – 1. use BERT embeddings and combine with non-transformer architectures, 2. take the pretrained BERT transformer model and fine tune by further training the model with toxic comment data.

D'Sa et al tried to apply three different pretrained word representations including BERT onto feature-based CNN and RNN models with an regression-based approach in [1]. A dataset similar to ours consists of 160,000 comments from the Wikipedia Detox project was used. The BERT fine-tuning approach achieved 78.2 F1-score, better than feature-based models that combined with BERT embeddings. The authors also tested robustness of the feature-based models by adding a toxic word like 'fuck' or a healthy word like 'love' to each comment of test set, and found out that model with BERT embedding is least susceptible to word appending attacks.

In Offensive Language Identification and Categorization with Perspective and BERT by Pavlopoulos, Androutsopoulos, Thain and Dixon published in 2019 [4], authors observed that Perspective(a CNN trained by Jigsaw and Google on millions of user comments from different online publishers based on GloVe embeddings) performed better at macro F1-score of 0.7933, than BERT fine-tuning(0.7705) in detecting toxicity. But in terms of subcategory detection tasks like threat identification, scores for both models are in general lower than toxicity identification. In this sense, BERT performed better in categorizing the offensive type(threats, insults, profanity, identity attack etc) with a 0.6817 F1-score, exceeding Perspective's 0.4785 by a wide margin.

4 DATASET

We used the dataset from the Kaggle Toxic Comment Classification Competition. The training set has 159,571 comments and every comment has 6 labels associated with it, namely toxic, severe toxic, obscene, threat, insult and identity hate. This is a multi-label dataset which means that a comment can have more than one class assigned to it. The assignments are represented by a 0 or 1 for each class. A 0 means the comment is not attributed to that label and 1 means that the comment is attributed to that specific label.

Even though there are 159,571 comments in total there is a high level of class imbalance and most of the sentences cannot be attributed to any of the classes which is why we have used down-sampling to deal with the problem. The total number of comments which has at least one kind of toxicity is approximately 15,000. So we have selected those comments and sampled another 15,000 comments where the comment cannot be attributed to any kind of toxicity. So in total we ended up with a dataset of around 30,000 comments. We have split that dataset into three parts - train set(20000 comments), validation set(20000 comments) and test set(5294 comments). The data is in primarily in English language and we worked with it in a CSV format.

5 METHODOLOGY

We have used Google Colab to run the experiments since none of us have access to GPU and it is not feasible to iterate on experiments that are run on a CPU. Since we know pretrained language models are quite efficient in doing a lot of downstream NLP tasks we have chosen to use BERT(Bidirectional Encoder Representations from Transformers) for our experiments. The Transformers library has been used to get these BERT Embeddings.

We have used the BertTokenizer to tokenize the comments and then add special tokens indicating the start and end of the sequence and also added padding to make the comment of a pre-determined length of 84 tokens. Our Neural Network consists of the BERT layer from Transformers and then a dropout layer with probability of 0.1. The dropout layer is followed by a Linear layer with output size of 6 because that is the total number of labels the comment can be attributed to.

We feed tokens to our Neural Network which outputs a vector of length 6. We then pass this output through a Sigmoid layer which squishes the values in between 0 and 1 and then we use BCELoss to calculate the combined loss of all the labels. We actually use the BCEWithLogitsLoss as the loss function because it combines both the sigmoid activation and BCELoss together and gives better numerical stability.

While making actual predictions we take the output from the Linear Layer and pass it through a Sigmoid function and then depending on whether the value at a particular position of the vector is above or below 0.5 we classify it as a 1 or 0 (>0.5 1 and <0.5 0). These ones and zeros essentially mean if the comment can be attributed to the label at that position.

To compare our results with a baseline model, we have trained a uni-directional LSTM model using the same processing steps and same loss functions and as expected the BERT model performs approximately 4% better than the baseline LSTM model. Our baseline model gives us a Macro-F1 Score of 0.86199 after training for 5 epochs. We have used a batch size of 32 for the training and the default learning rate of 0.001 using the Adam optimizer.

We used Macro-average version of F1-score due to imbalanced class nature of data, as the micro-average F1 is too lenient in giving out high scores even when some minority classes score low marks.

6 EXPERIMENTS

At first, we tried to train 6 binary classifiers, one for toxicity, and five for the remaining toxicity sub-categories, using the full train.csv which contains 159,571 comments. However, training merely one epoch with a lower-memory-consuming bert-base-uncased embedding already took an hour to finish, though the validation result on dev-set is quite good at a macro F1-score of 0.8844.

It would be too time-consuming to train the rest 5 binary classifiers with all the 159,571 comments, and a workaround intended was to train the subcategory classifiers only on the 15,294 comments that were annotated to be toxic, yet the results are much worse for subcategories

Table 1: Fine tuning learning rate with different epochs of Macro-F1 score for test set

Epochs+Learning rate	5e-05(Batch size 16)	3e-05(Batch size 16)	2e-05(Batch size 16)	2e-05(Batch size 32)
2 Epochs	90.40%	89.86%	89.96%	90.65%
3 Epochs	90.04%	89.59%	89.66%	90.38%
4 Epochs	89.89%	89.78%	89.41%	90.33%

Table 2: Fine tuning max_grad_norm with warm_up ratio of Macro-F1 score for test set

max_grad_norm + warm_up ratio	0.05	0.1	0.15	0.2
max_grad_norm at 0.7	90.50%	90.32%	90.15%	89.91%
max_grad_norm at 0.8	90.09%	89.93%	89.99%	89.79%
max_grad_norm at 0.9	90.12%	89.85%	89.83%	89.95%
max_grad_norm at 1.0	90.07%	90.09%	90.01%	89.97%
max_grad_norm at 1.1	89.98%	89.79%	89.79%	89.87%

that have very few positive labels, due to class imbalance problem. For example, the "identity_hate" column only has 1,405 positive examples among all 15,294 toxic comments, and therefore F1-score on dev-set was bad at 0.4776 using 'bert-base-uncased' embedding, while much better at 0.7981 using 'bert-large-uncased' embedding(perhaps addressing out-of-vocabulary problem). But due to Colab memory constraint, we cannot train all classifiers with 'bert-large-uncased'.

Some attempts to train other subcategories including insult, obscene and severe_toxic achieved validation F1-score between 0.7132 and 0.8159 after 3 epochs of training, which are far worse than the training result on toxicity that were trained on full data.

Afterwards, we switched to a better architecture that can train 1 epoch in just 3-6 minutes, and can train 6 labels for each sentence in-one-go. The new model achieved around 0.90 macro-F1 score on our test split that contained about 5,000 sentences which was sampled from original training dataset.

And then we tried finetuning hyperparameters in the following ranges according to suggestions in the original BERT paper(see below):

1. Batch size: 16, 32
2. Learning rate (Adam): 5e-5, 3e-5, 2e-5
3. Number of epochs: 2, 3, 4

Finetuning results with 6-labels-in-one-go architecture are surprising very similar, always hovering between 0.89 - 0.90, with best combination of (Batch_size 32, Num_epochs 2 and learning rate 2e-5) at 0.9065. The detailed results are showed in table 1.

Further hyperparameter tuning on max_grad_norm and warmup_proportion also produces test-split macro-average F1 between 0.89 - 0.905, with best combination of (max_grad_norm at 0.7 and warmup_proportion at 0.05), building on the base combination of (Batch_size 32, Num_epochs 2 and learning rate 2e-5). This could mean our model is quite robust. Table2 shows some minor variations(highest score is not as good as 0.9065 this time probably due to some randomness in model runs).

7 RESULTS

We tested the BERT model performance with the Kaggle test data which contains 153,164 comments. Our model reports a lower Macro F1 score at 0.788, lower than the nice >0.90 scores we had on the test split of training data. We submitted the predictions to the Kaggle competition which scores differently. It takes in probabilities(sigmoid values) and not the actual 0 and 1 classification labels for the six columns. The evaluation metric is mean column-wise ROC AUC, according to the competition page, which measures how accurate model labelings are for each of the 6 categories and then take an average. The graded score is 0.97905 and such a score could rank about 2586th on the public leaderboard among 4500+ entries (The competition is closed already).


We prepared 5,494 test sentences in total to further evaluate the performance of our model. The classifier is pretty good based on the F1 score and it reaches 0.90 in our test split. We processed 10 epochs to train our model and get the best performance on epoch 9. In order to find the rest of errors an error analysis is performed based on the incorrect predictions. Precision score for non_toxic (negative class labelled as '0') is 0.96, which is higher than precision for toxic class(positive class labelled as "1") under six categories which refer to toxic, severe_toxic, obscene, threat, insult and identity_hate.

Besides, based on the classification report below, we can see that precision for non-toxic class is 0.961 and 0.854 for toxic class. The recall is 0.969 for non-toxic but 0.820 for toxic comments. The difference between precision of 0.961 and recall of 0.969 for non-toxic class is very small. That is to say, our model has better performance on predicting non-toxic comments rather than toxic comments. Please refer to more detailed statistical results in the table 3.

Compared with our gold label, there are 1,398 comments in our test split which have the incorrect prediction labels. One of the reason for false predictions is multiple labels for one sentence. For example, "the stupid one, not me". This one is labeled as toxic under both toxic and obscene categories based on our actual labels, however it is predicted as only toxic. In this situation, our model is more accurate apparently,

Table 3: Confusion matrix for test set

	precision	recall	f1score	support
non_toxic	96.1%	96.9%	96.5%	26010
toxic	85.4%	82.0%	83.7%	5754
accuracy			94.2	31764
macro avg	90.7%	89.5%	90.1%	31764
weighted avg	94.1%	94.2%	94.2%	31764

**Figure 1: Visualization for attention analysis of a toxic sentence**

since there is no obscene words in this sentence. Another reason is that we have wrong gold labels. Let's take one sentence as an example. "im in your area i ' m going to find you and when i find you i will teach you how to swim." Our classifier predict the sentence as nontoxic for all six categories, but the gold standard is labeled as toxic under category toxic. Obviously this sentence is nontoxic in any aspect. Gold labels are not fully accurate since they are labeled by human raters for toxic behavior.

Besides, we take one sentence as an example to process attention analysis. For example, a sentence like "stop fucking with the y tm nd wiki" from our dataset is toxic apparently. And we have the higher attention weight in head of 3 and head of 12 for the bad word of the second token, which shows our model is pretty good to predict the label as toxic.

Based on the attention analysis visualization, as we can observe that X-axis is the word tokens/key which attention is being paid, and Y-axis is number of head. The intensity of color blue shows attention weights. As we can observe, the bad words in this sentence have more attention weights higher than 0.35, which can be a good indicator to predict the comment as toxic. General speaking, our model forms a strong composite representations to understand language.

8 CONCLUSION

To conclude, the best combination of our model is batch size of 32, epochs of 2 and learning rate of $2e-5$ which we obtained F1-score of 0.905. Our main BERT model performs approximately 4% better than the baseline uni-directional LSTM model.

We encountered some challenges. One of the challenges we faced is that this is not a typical binary or a simple multi-class classification problem where every comment can only be attributed to one class. A comment can be attributed to more than one class which meant we cannot use softmax in a typical way. The immediate solution was to have 6 different models where every model will be a binary classifier for one class predicting if the comment can or cannot be attributed to that specific class. However, this was very inefficient from a computation standpoint since we would have to use 6 times the memory and storage compared to a single model and also spend way more time iterating

on each of the models. The more non-trivial solution we came up with is to use a single network for all classifying all the labels and use the BCELoss along with the sigmoid activation function to predict and calculate the loss in one go.

In the future, we will study more on dealing with data imbalance and possibly figure out a way to add all the extra data we did not use without making the model biased towards non-toxic comments. We would also like to extend our model to apply XLNet pretrained model as well.

REFERENCES

- [1] D'Sa, A. G., Illina, I., and Fohr, D. (2019). Towards non-toxic landscapes: Automatic toxic comment detection using dnn. *arXiv preprint arXiv:1911.08395*.
- [2] Google (2017). Attention is all you need. *arXiv preprint arXiv:1706.03762*.
- [3] Khieu, K. and Narwal, N. (2017). Detecting and classifying toxic comments.
- [4] Pavlopoulos, J., Thain, N., Dixon, L., and Androutsopoulos, I. (2019). Convai at semeval-2019 task 6: Offensive language identification and categorization with perspective and bert. In *Proceedings of the 13th International Workshop on Semantic Evaluation*, pages 571–576.
- [5] van Aken, B., Risch, J., Krestel, R., and Löser, A. (2018). Challenges for toxic comment classification: An in-depth error analysis. *arXiv preprint arXiv:1809.07572*.